



SegmentationFault Fest

# ALBERTO VARA



# DESCARGAR PRESENTACIÓN/CODIGO FUENTE

[github.com/avara1986/python-talks/2022\\_10\\_pycones](https://github.com/avara1986/python-talks/2022_10_pycones)

# MOTIVACION

# MOTIVACION



# ESCENARIO

- Instrumentar código
- Reducir al mínimo el tiempo extra (overhead)
- Python llega a un punto que no es óptimo

# ESCENARIO

- Instrumentar código
  - "Wrappear" código
  - Monkey Patching
- Reducir al mínimo el tiempo extra (overhead)
- Python llega a un punto que no es óptimo

func\_1 -> 3.06 µs. func\_6 -> 5.43 µs

```
def func_1(a, b, c):
    return f"hello {a} {b} {c}"

def func_2(a, b, c):
    return func_1(a, b, c)

def func_3(a, b, c):
    return func_2(a, b, c)

def func_4(a, b, c):
    return func_3(a, b, c)

def func_5(a, b, c):
    return func_4(a, b, c)
```

# HERRAMIENTAS/LIBRERÍAS

- Cython
- Pyperf
- Py-spy
- Gdb

# EN MATERIA

- Rendimiento
- Segmentation Faults

# PERFORMANCE



# BÚSQUEDA BINARIA PYTHON

```
# binarysearch.py
def binary_search(l: List[int], value: int, low: int, high: int):
    if high >= low:
        mid = (high + low) // 2
        if l[mid] == value:
            return mid
        elif l[mid] > value:
            return binary_search(l, value, low, mid - 1)
        else:
            return binary_search(l, value, mid + 1, high)
    else:
        return -1
```

# BÚSQUEDA BINARIA PYTHON

```
SEARCH_LIST = [i for i in range(100000)]
```

```
# binarysearch.py
def benchmark_binary_search(loops: int = 10) -> None:
    len_search_list = len(SEARCH_LIST)
    for i in range(0, loops):
        binary_search(SEARCH_LIST, 66666, 0, len_search_list)
```

# BÚSQUEDA BINARIA CYTHON

```
# _binarysearch.pyx
cdef int cython_binary_search(l: List[int], value: int, low: int, high
    if high >= low:
        mid = (high + low) // 2
        if l[mid] == value:
            return mid
        elif l[mid] > value:
            return cython_binary_search(l, value, low, mid - 1)
        else:
            return cython_binary_search(l, value, mid + 1, high)
    else:
        return -1
```

# BÚSQUEDA BINARIA C++

```
# src/binarysearch.hpp
int binsearch(vector<int> my_list, int value, int low, int high) {
    if (high >= low) {
        int mid = (low + high) / 2;
        if (my_list[mid] == value) {
            return mid;
        } else if (my_list[mid] > value) {
            return binsearch(my_list, value, low, mid - 1);
        } else {
            return binsearch(my_list, value, mid + 1, high);
        }
    } else {
        return -1;
    }
}
```

# BÚSQUEDA BINARIA C++

```
# _binarysearch.pyx
cdef extern from "src/binarysearch.hpp":
    int binsearch(vector[int] l, int value, int low, int high)

cpdef void benchmark_native_binary_search(loops=10):
    cdef int len_search_list = len(SEARCH_LIST)
    for i in range(0, loops):
        binsearch(SEARCH_LIST, 66666, 0, len_search_list)
```

# BENCHMARKS CON PYPERF

# BENCHMARKS CON PYPERF

```
python benchmark.py -f native_binary_search -i 10 -o native_binary_sear
python benchmark.py -f cython_binary_search -i 10 -o cython_binary_sear
python benchmark.py -f binary_search -i 10 -o binary_search_10.json

python -m pyperf compare_to --table binary_search_10.json \
    cython_binary_search_10.json \
    native_binary_search_10.json
```

Sin optimización. 10 bucles

Python	Cython	C++
30.9 us	21.2 us: 1.46x faster	42.5 ms: 1372.82x slower



```
#include <chrono>
#include "binarysearch.hpp"

int main(int argc, char *argv[]) {
    vector<int> arr;
    for (int i = 0; i < 1000000; i++) arr.push_back(i);

    std::chrono::steady_clock::time_point begin = std::chrono::steady_c
    for (int i = 0; i < 10; i++) {
        binsearch(arr, 66666, 0, arr.size());
    }
    std::chrono::steady_clock::time_point end = std::chrono::steady_clo
    cout << "Elapsed time: " << std::chrono::duration<double>(end - begin)
```

```
cmake . && make && ./binarysearch
```

```
./binarysearch
Elapsed time: 2.994[ms]
Elapsed time: 2.994.433[μs][microseconds]
Elapsed time: 2.994.433.429[ns][nanoseconds]
```

# BÚSQUEDA BINARIA C++

```
1 int binsearch(vector<int> my_list, int value, int low, int high) {  
2     if (high >= low) {  
3         int mid = (low + high) / 2;  
4         if (my_list[mid] == value) {  
5             return mid;  
6         } else if (my_list[mid] > value) {  
7             return binsearch(my_list, value, low, mid - 1);  
8         } else {  
9             return binsearch(my_list, value, mid + 1, high);  
10        }  
11    } else {  
12        return -1;  
13    }  
14 }
```

# BÚSQUEDA BINARIA C++

```
1 int binsearch(vector<int>& my_list, int value, int low, int high) {  
2     if (high >= low) {  
3         int mid = (low + high) / 2;  
4         if (my_list[mid] == value) {  
5             return mid;  
6         } else if (my_list[mid] > value) {  
7             return binsearch(my_list, value, low, mid - 1);  
8         } else {  
9             return binsearch(my_list, value, mid + 1, high);  
10        }  
11    } else {  
12        return -1;  
13    }  
14 }
```

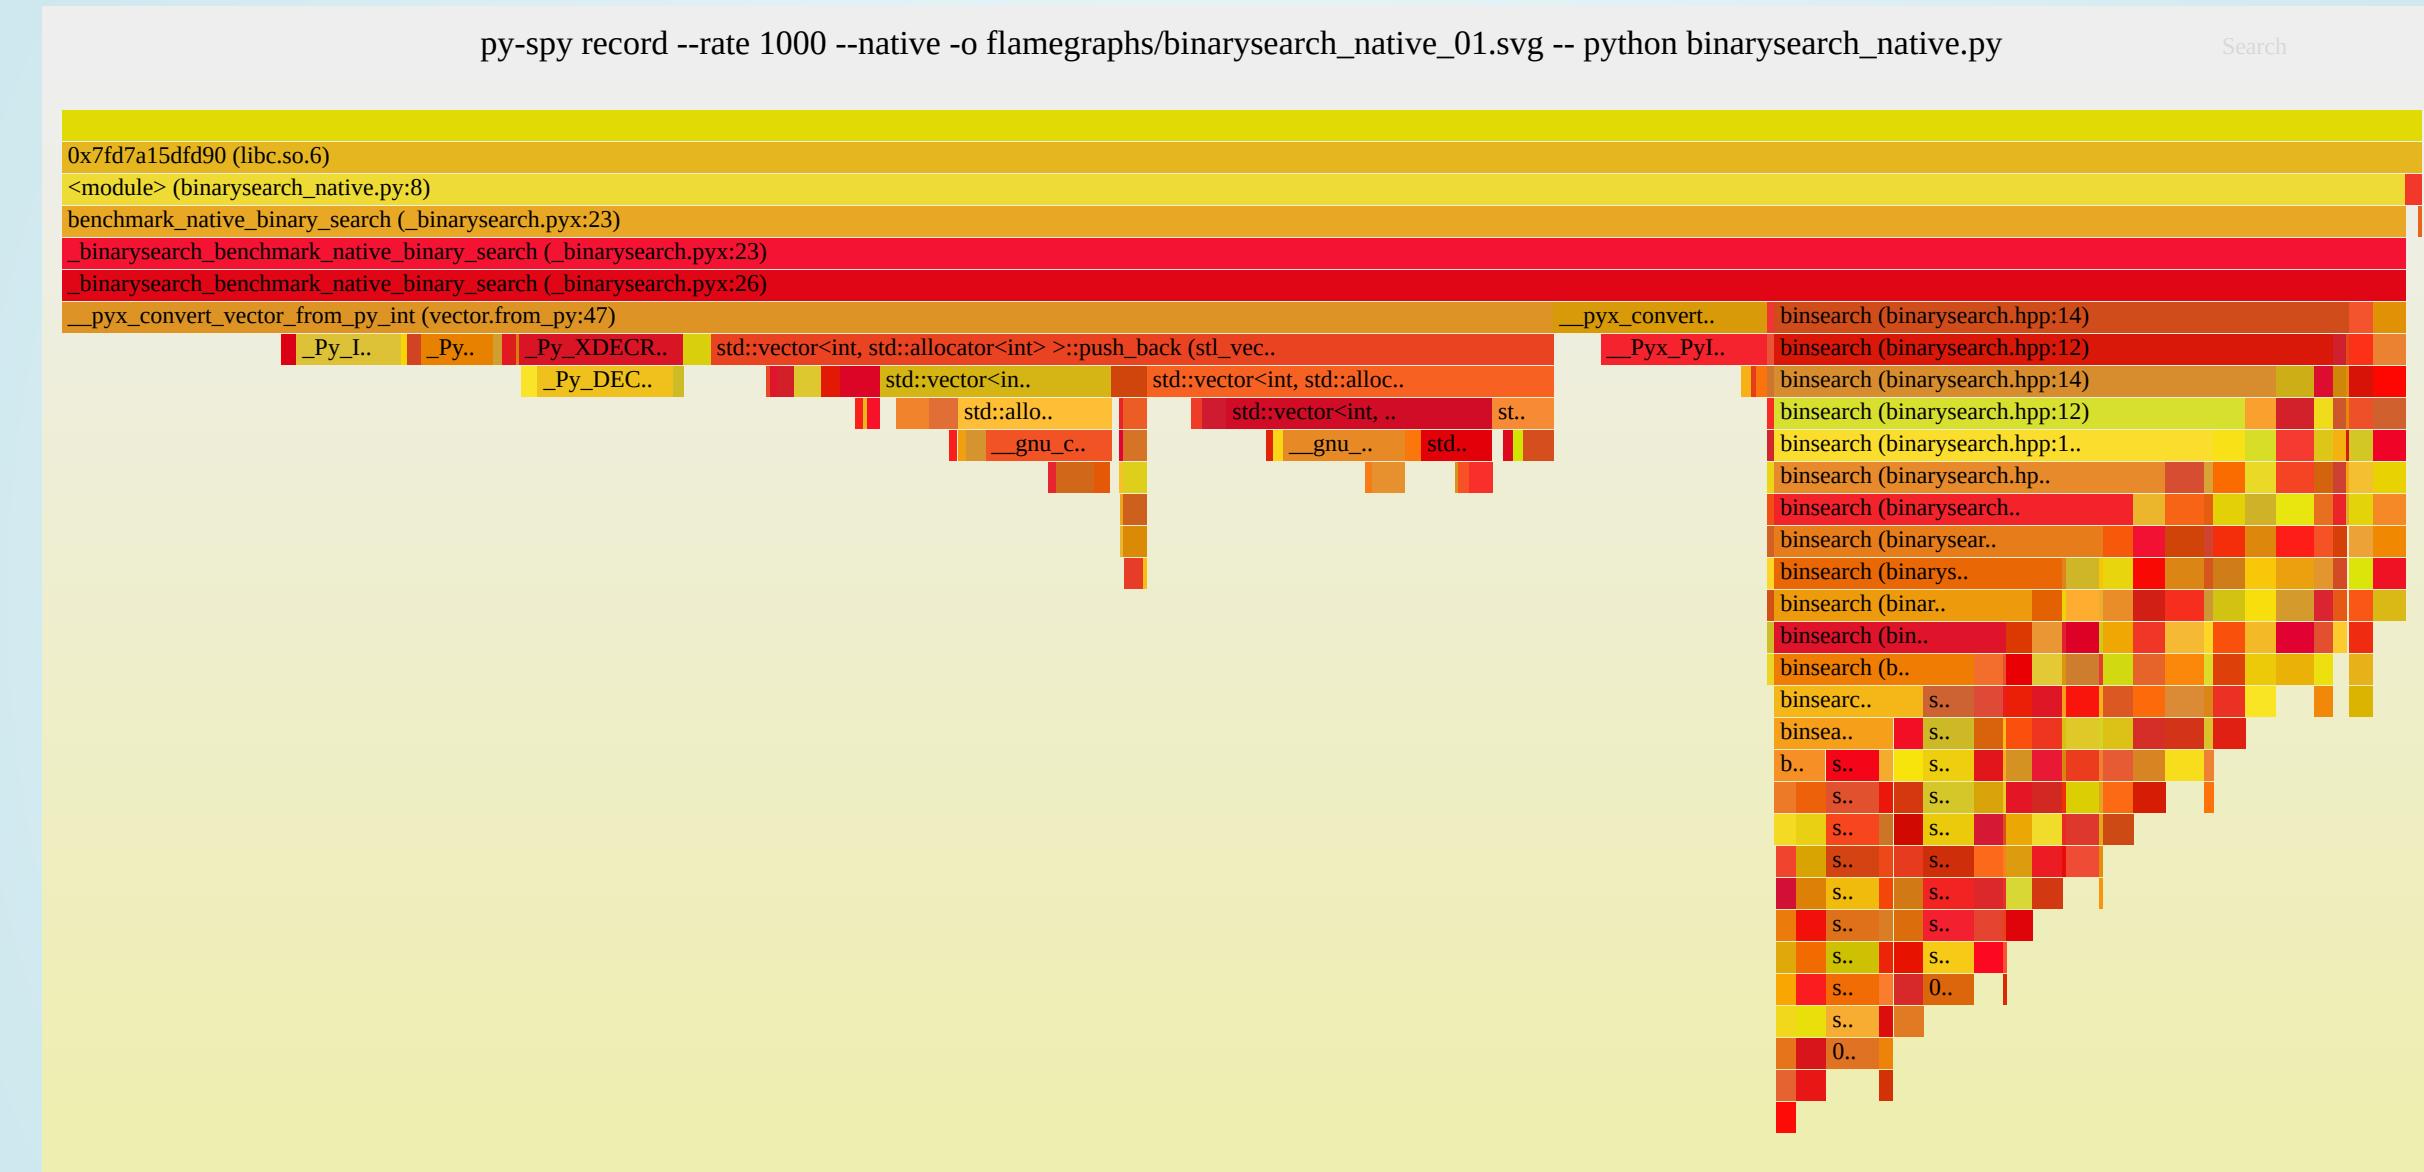
```
cmake . && make && ./binarysearch
```

```
./binarysearch
Elapsed time: 0[ms]
Elapsed time: 2[μs][microseconds]
Elapsed time: 2.827[ns][nanoseconds]
```

Con optimización. 10 bucles

Python	Cython	C++
30.6 us	20.8 us: 1.47x faster	24.2 ms: 792.40x slower

```
py-spy record --rate 1000 --native \  
  -o flamegraphs/binarysearch_native_01.svg -- python binarysearch_na
```



```
1 cpdef void benchmark_native_binary_search(loops=10):
2     cdef int len_search_list = len(SEARCH_LIST)
3     for i in range(0, loops):
4         binsearch(SEARCH_LIST, 66666, 0, len_search_list)
```

```
1 cpdef void benchmark_native_binary_search(loops=10):
2     cdef int len_search_list = len(SEARCH_LIST)
3     cdef vector[int] my_list = SEARCH_LIST
4     for i in range(0, loops):
5         binsearch(&my_list, 66666, 0, len_search_list)
```

# BÚSQUEDA BINARIA C++

```
1 int binsearch(vector<int>& my_list, int value, int low, int high) {  
2     if (high >= low) {  
3         int mid = (low + high) / 2;  
4         if (my_list[mid] == value) {  
5             return mid;  
6         } else if (my_list[mid] > value) {  
7             return binsearch(my_list, value, low, mid - 1);  
8         } else {  
9             return binsearch(my_list, value, mid + 1, high);  
10        }  
11    } else {  
12        return -1;  
13    }  
14 }
```

# BÚSQUEDA BINARIA C++

```
1 int binsearch(vector<int>* my_list, int value, int low, int high) {  
2     if (high >= low) {  
3         int mid = (low + high) / 2;  
4         if ((*my_list)[mid] == value) {  
5             return mid;  
6         } else if ((*my_list)[mid] > value) {  
7             return binsearch(my_list, value, low, mid - 1);  
8         } else {  
9             return binsearch(my_list, value, mid + 1, high);  
10        }  
11    } else {  
12        return -1;  
13    }  
14 }
```

Otimización con punteros. 10 bucles

Python	Cython	C++
31.6 us	21.0 us: 1.51x faster	2.55 ms: 80.80x slower



Y si...

## Otimización con punteros. 1.000 bucles

Python	Cython	C++
3.06 ms	1.97 ms: 1.55x faster	2.66 ms: 1.15x faster



## Optimización con punteros. 10.000 bucles

Python	Cython	C++
30.2 ms	19.5 ms: 1.55x faster	3.91 ms: 7.73x faster

Optimización con punteros. 100.000 bucles

Python	Cython	C++
304 ms	193 ms: 1.58x faster	14.9 ms: 20.40x faster



# SEGMENTATION FAULT (FEST!)



```
strings_and_bytes_dict = {  
    "barba": "coa",  
    "aaa": b"value_aaa",  
    "bbb": b"value_bbb",  
    "ccc": "value_ccc",  
}  
bytes_dict = BytesDict(strings_and_bytes_dict)  
print(bytes_dict)
```

```
cdef class BytesDict(object):
    cdef dict bytes_dict

    def __init__(self, dict strings_dict):
        self.bytes_dict = {}
        cdef char * ptr
        cdef char * bytes_key = NULL
        cdef char * bytes_value = NULL

        for k, v in strings_dict.items():
            k = self._string_to_bytes(k, &ptr)
            bytes_key = ptr
            v = self._string_to_bytes(v, &ptr)
            bytes_value = ptr
```

```
cdef object _string_to_bytes(self, object string, char ** ptr):
    if isinstance(string, six.binary_type):
        ptr[0] = PyBytes_AsString(string)
    elif isinstance(string, six.text_type):
        if ptr[0] == NULL:
            string = PyUnicode_AsEncodedString(string, "utf-8", "surrogates")
            ptr[0] = PyBytes_AsString(string)
    return string
```

```
python setup.py build_ext --inplace
python segmentation.py
[1]    40755 segmentation fault (core dumped)  python segmentation.py
```

```
gdb --args python segmentation.py
Reading symbols from python...
(No debugging symbols found in python)
(gdb)
```

```
(gdb) r
Starting program: /home/alberto.vara/mio/python-talks/2022_10_pycones/
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so

Program received signal SIGSEGV, Segmentation fault.
0x000000000006b842e in PyBytes_FromString ()
gdb)
```

```
1 gdb) backtrace
2 #0 0x000000000006b842e in PyBytes_FromString ()
3 #1 0x00007ffff7fb4b48 in __pyx_f_13_segmentation_9BytesDict__update_dict (
4 #2 0x00007ffff7fb4918 in __pyx_pf_13_segmentation_9BytesDict__init__ (__p
5 #3 0x00007ffff7fb3fa0 in __pyx_pw_13_segmentation_9BytesDict_1__init__ (__p
6 #4 0x00000000005f745f in _PyObject_MakeTpCall ()
7 #5 0x0000000000570d55 in _PyEval_EvalFrameDefault ()
8 #6 0x0000000000569dba in _PyEval_EvalCodeWithName ()
9 #7 0x00000000006902a7 in PyEval_EvalCode ()
10 #8 0x000000000067f951 in ?? ()
11 #9 0x000000000067f9cf in ?? ()
12 #10 0x000000000067fa71 in ?? ()
13 #11 0x0000000000681b97 in PyRun_SimpleFileExFlags ()
14 #12 0x0000000000000000 in PyRunMain (...)
```

```
1 gdb) backtrace
2 #0 0x000000000006b842e in PyBytes_FromString ()
3 #1 0x00007ffff7fb4b48 in __pyx_f_13_segmentation_9BytesDict__update_dict (
4 #2 0x00007ffff7fb4918 in __pyx_pf_13_segmentation_9BytesDict__init__ (__p
5 #3 0x00007ffff7fb3fa0 in __pyx_pw_13_segmentation_9BytesDict_1__init__ (__p
6 #4 0x00000000005f745f in _PyObject_MakeTpCall ()
7 #5 0x0000000000570d55 in _PyEval_EvalFrameDefault ()
8 #6 0x0000000000569dba in _PyEval_EvalCodeWithName ()
9 #7 0x00000000006902a7 in PyEval_EvalCode ()
10 #8 0x000000000067f951 in ?? ()
11 #9 0x000000000067f9cf in ?? ()
12 #10 0x000000000067fa71 in ?? ()
13 #11 0x0000000000681b97 in PyRun_SimpleFileExFlags ()
14 #12 0x0000000000000000 in PyRunMain ()
```

```
1 gdb) backtrace
2 #0 0x00000000006b842e in PyBytes_FromString ()
3 #1 0x00007ffff7fb4b48 in __pyx_f_13_segmentation_9BytesDict__update_dict (
4 #2 0x00007ffff7fb4918 in __pyx_pf_13_segmentation_9BytesDict__init__ (__p
5 #3 0x00007ffff7fb3fa0 in __pyx_pw_13_segmentation_9BytesDict_1__init__ (__p
6 #4 0x00000000005f745f in _PyObject_MakeTpCall ()
7 #5 0x0000000000570d55 in _PyEval_EvalFrameDefault ()
8 #6 0x0000000000569dba in _PyEval_EvalCodeWithName ()
9 #7 0x00000000006902a7 in PyEval_EvalCode ()
10 #8 0x000000000067f951 in ?? ()
11 #9 0x000000000067f9cf in ?? ()
12 #10 0x000000000067fa71 in ?? ()
13 #11 0x0000000000681b97 in PyRun_SimpleFileExFlags ()
14 #12 0x0000000000000000 in PyRunMain (...)
```

```
1 for k, v in strings_dict.items():
2     k = self._string_to_bytes(k, &ptr)
3     bytes_key = ptr
4     v = self._string_to_bytes(v, &ptr)
5     bytes_value = ptr
6     self._update_dict(bytes_key, bytes_value)
```

```
1 cdef object _string_to_bytes(self, object string, char ** ptr):
2     if isinstance(string, six.binary_type):
3         ptr[0] = PyBytes_AsString(string)
4     elif isinstance(string, six.text_type):
5         if ptr[0] == NULL:
6             string = PyUnicode_AsEncodedString(string, "utf-8", "surrogatepa
7             ptr[0] = PyBytes_AsString(string)
8     return string
```

```
1 cdef object _string_to_bytes(self, object string, char ** ptr):
2     ptr[0] = NULL
3     if isinstance(string, six.binary_type):
4         ptr[0] = PyBytes_AsString(string)
5     elif isinstance(string, six.text_type):
6         if ptr[0] == NULL:
7             string = PyUnicode_AsEncodedString(string, "utf-8", "surrogatepa
8             ptr[0] = PyBytes_AsString(string)
9     return string
```

```
1  def __init__(self, dict strings_dict):
2      self.bytes_dict = {}
3      cdef char * ptr
4      cdef char * bytes_key = NULL
5      cdef char * bytes_value = NULL
6
7      for k, v in strings_dict.items():
8          k = self._string_to_bytes(k, &ptr)
9          bytes_key = ptr
```

```
1  def __init__(self, dict strings_dict):
2      self.bytes_dict = {}
3      cdef char * ptr = NULL
4      cdef char * bytes_key = NULL
5      cdef char * bytes_value = NULL
6
7      for k, v in strings_dict.items():
8          k = self._string_to_bytes(k, &ptr)
9          bytes_key = ptr
```

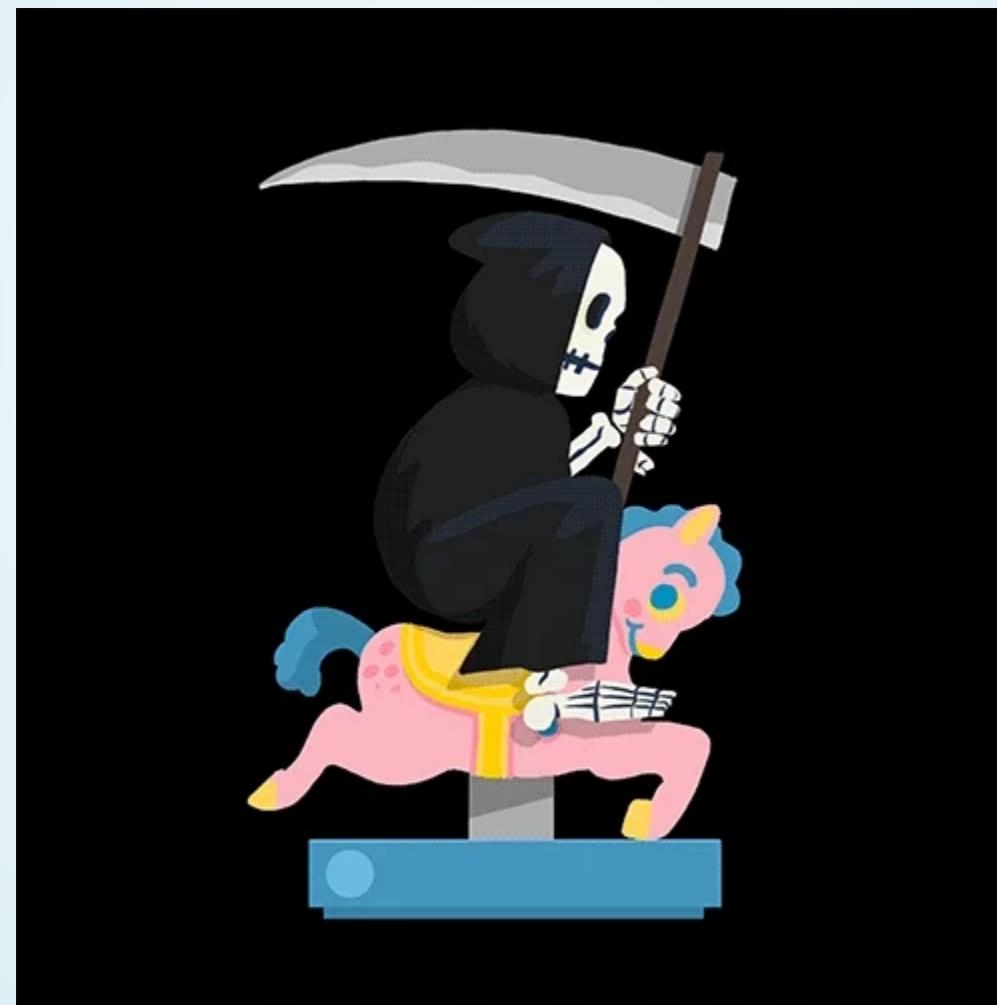
```
python setup.py build_ext --inplace  
python segmentation.py
```

```
{b'barba': b'value_aaa', b'value_aaa': b'value_bbb', b'value_bbb': b'velho', b'velho': b'barba'}
```

# CONCLUSIONES

# PASAR VARIABLES ENTRE LENGUAJES TIENE UN PEAJE

P.E: Como la latencia de red entre microservicios



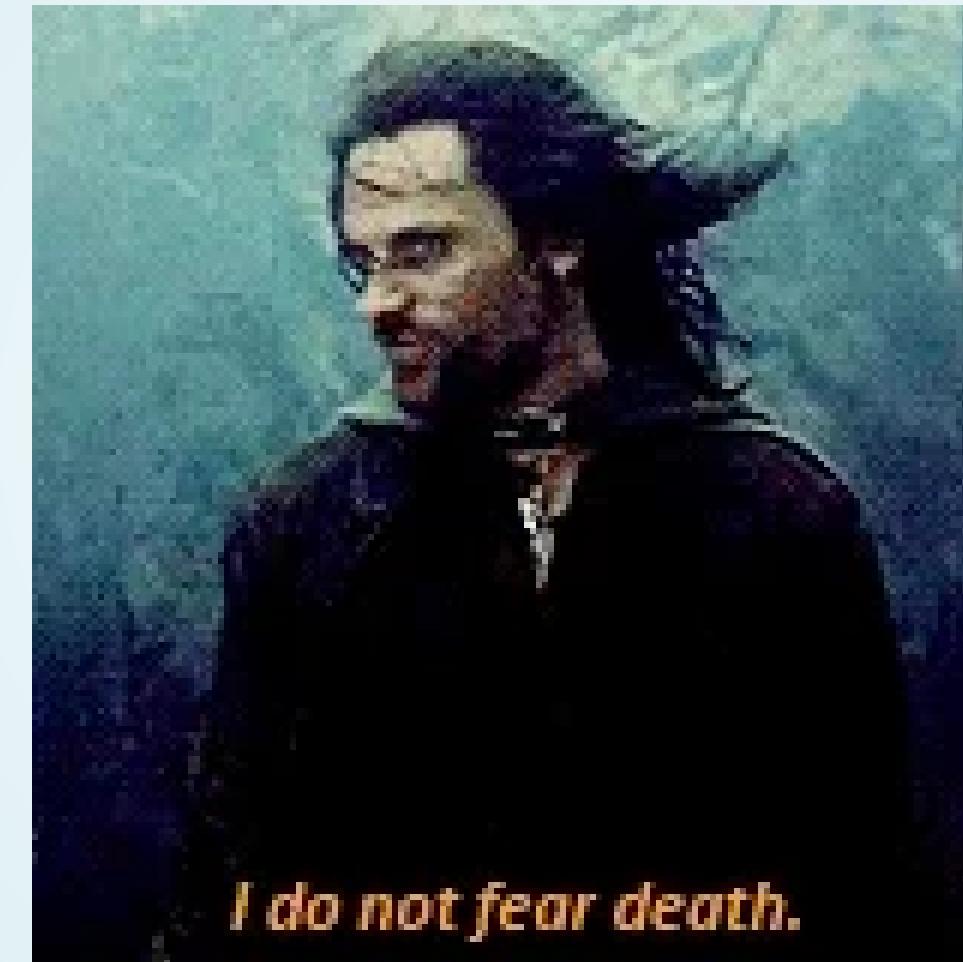
# NO TODO SE PUEDE MIGRAR

Para compensar el peaje, el numero de operaciones tiene que ser considerable, lo que dificulta una migración

# CAMBIO DE PARADIGMA

IMHO, Cambio de paradigma lenguajes como Python, PHP, JS a compilados C o C++ es mayor

# ¿HE MENCIONADO EL USO DE MEMORIA EN C++?



¡MUCHAS GRACIAS POR VUESTRO TIEMPO!



We are hiring!  
[www.datadoghq.com/careers/](http://www.datadoghq.com/careers/)